



# Ozone Analysis



Leon Zha  
ITP449  
James Faghmous



# EPA Air Quality Monitor Dataset

---

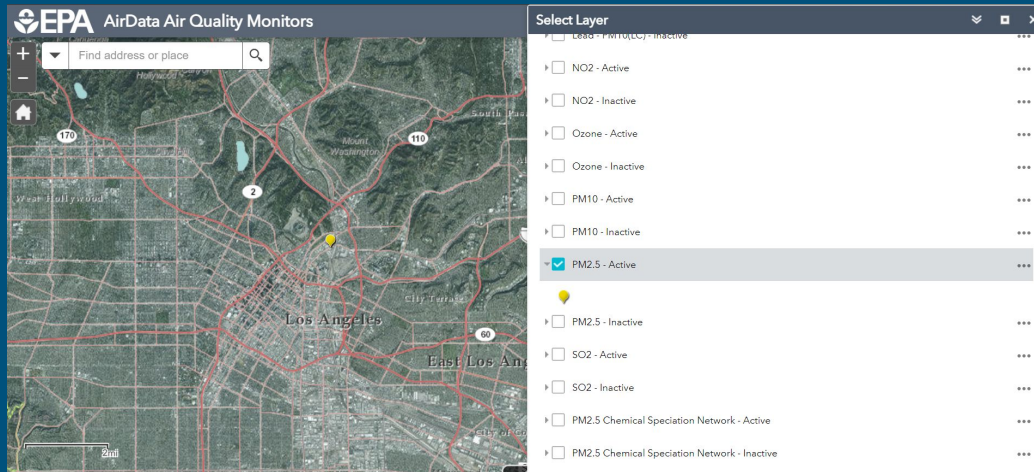
- Somewhat of a continuation of my Exploratory Data Analysis project
- Same original dataset, different subset
  - This time, I'm looking at ozone

# What is ozone?

---

- Ozone is an unstable gas comprised of three oxygen atoms
- Vital to our survival in the upper atmosphere (stratosphere) because it absorbs UV radiation from the Sun
- However, in the lower atmosphere where we live (troposphere) ozone is a pollutant that can harm us
  - Damages our cells, causing asthma attacks, eye and lung irritation, heart disease, etc.
  - Also damages materials such as rubber, paints, fibers, etc.
- Tropospheric ozone formation relies on UV light from the sun, nitrogen oxides, and hydrocarbons
- Source: <https://scied.ucar.edu/learning-zone/air-quality/ozone-troposphere>

# Data Source



- <https://epa.maps.arcgis.com/apps/webappviewer/index.html?id=5f239fd3e72f424f98ef3d5def547eb5&extent=-116.2334,113.1913,-46.3896,56.5319>
- Los Angeles-North Main Street Station Monitor
- Looking at PM2.5 Chemical Speciation Network - Active
- Using daily data from 2016-2020

# What am I doing in this notebook?

---

- Goal: predict ozone levels around LA based on other variables such as temperature, nitrogen dioxide, nitrogen oxide, etc.
- I will use linear regression and ensemble methods


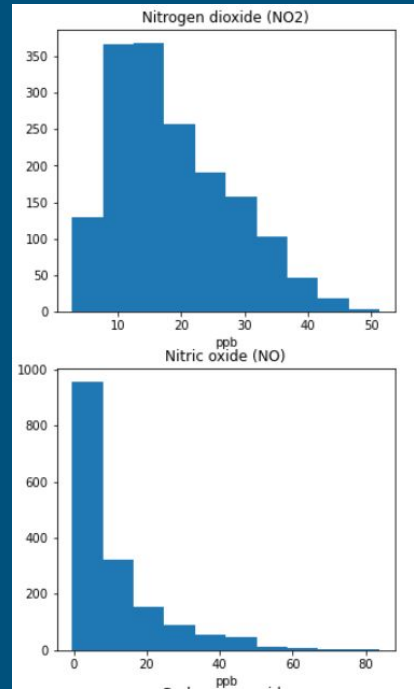
# After reading in and cleaning the data...

- We see there are 1640 rows and 6 columns
- No null values (dropped)
- A look at part of the data

```
<bound method DataFrame.info of Parameter    Average Ambient Temperature    Carbon monoxide    Nitric oxide (NO) \
Date
2016-01-01          11.82          0.441777          4.970019
2016-01-02          14.00          0.546181          16.287682
2016-01-03          13.50          0.787658          29.230435
2016-01-04          16.20          0.444372          8.648642
2016-01-05          14.60          0.362291          7.527083
...
2020-09-26          22.60          0.347917          2.002084
2020-09-27          22.80          0.297916          0.964765
2020-09-28          24.20          0.425000          4.706250
2020-09-29          26.30          0.522917          8.384511
2020-09-30          29.40          0.654167          19.339583
```



# Looking at distributions

- Next, I created histograms of each variable to look at their distributions
- Most of them had tails to the right, but others were distributed more evenly
- Only four of these are shown to the right



# More exploratory analysis

- Next, I graphed ozone levels against each of the other variables
- Some, like nitrogen dioxide, exhibit negative relationships
  - Probably because the nitrogen dioxide is being converted into ozone
- Temperature, on the other hand, exhibits a positive relationship
  - More temperature probably means it's a sunnier day, meaning more UV light to power ozone formation





# Attempt #1

- Since ozone levels are continuous, I went for a simple linear regression model
- I split the training and testing sets 70/30, using random\_state=2020 to ensure consistency between trials
- The base model, when tested against the testing data, resulted in a score of 0.5731489871678459
- The score is derived from scikit-learn's score() method, which calculates the difference between the model's predicted values and the actual predicted values (in other words, accuracy)

```
#now, getting into the actual machine learning part

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn import preprocessing

x_train, x_test, y_train, y_test = train_test_split(variables, target, test_size=0.3, random_state=2020)

#Something to note is that the axes from the plots earlier
#vary a lot (ppb vs ppm)
#To account for this let's scale
x_train = preprocessing.scale(x_train)
x_test = preprocessing.scale(x_test)
y_train = preprocessing.scale(y_train)
y_test = preprocessing.scale(y_test)

lin_model = LinearRegression()
lin_model.fit(x_train,y_train)

test_pred = lin_model.predict(x_test)

print("Training data score:", lin_model.score(x_train,y_train))
print("Testing data score:", lin_model.score(x_test,y_test))

Training data score: 0.5804274368655039
Testing data score: 0.5731489871678459
```

# Making a better model

- To increase the performance of the base model I tried simplifying it through cross validation
- I chose selected the optimal alpha by generating a list of potentials and finding the one with the best performance
- First, I used lasso, which resulted in a slightly better score of 0.5744636661133027

```
#trying to improve the model with cross validation
from sklearn.linear_model import Lasso
from sklearn.linear_model import Ridge
from sklearn.model_selection import cross_val_score

alpha_space = np.logspace(-4, 0, 50)

model_scores = []

lasso_model = Lasso(normalize=True)
for alpha in alpha_space:
    |
    # Specify the alpha value to use
    lasso_model.alpha = alpha

    # Perform 10-fold CV
    lasso_cv_scores = cross_val_score(lasso_model,x_train,y_train,cv=10)

    # Append the mean of lasso_cv_scores to model_scores = []
    model_scores.append(np.mean(lasso_cv_scores))

print("Training data score using lasso:", np.max(model_scores))

bestAlphaIndex = np.argmax(model_scores)
lasso_model.alpha = alpha_space[bestAlphaIndex]
lasso_model.fit(x_train, y_train)
print("Testing data score using lasso: ", lasso_model.score(x_test, y_test))

Training data score using lasso: 0.45935495585064673
Testing data score using lasso: 0.48284201892494427
```

# Making a better model cont.

- Next, I tried using ridge
- Similarly to lasso, I generated a range of alpha values to find the optimal one
- Ridge performed slightly worse than lasso on the testing set, with a score of 0.5737411667807468

```
model_scores = []

ridge_model = Ridge(normalize=True)
for alpha in alpha_space:

    # Specify the alpha value to use
    ridge_model.alpha = alpha

    # Perform 10-fold CV
    ridge_cv_scores = cross_val_score(ridge_model, x_train, y_train, cv=10)

    # Append the mean of lasso_cv_scores to model_scores = []
    model_scores.append(np.mean(ridge_cv_scores))

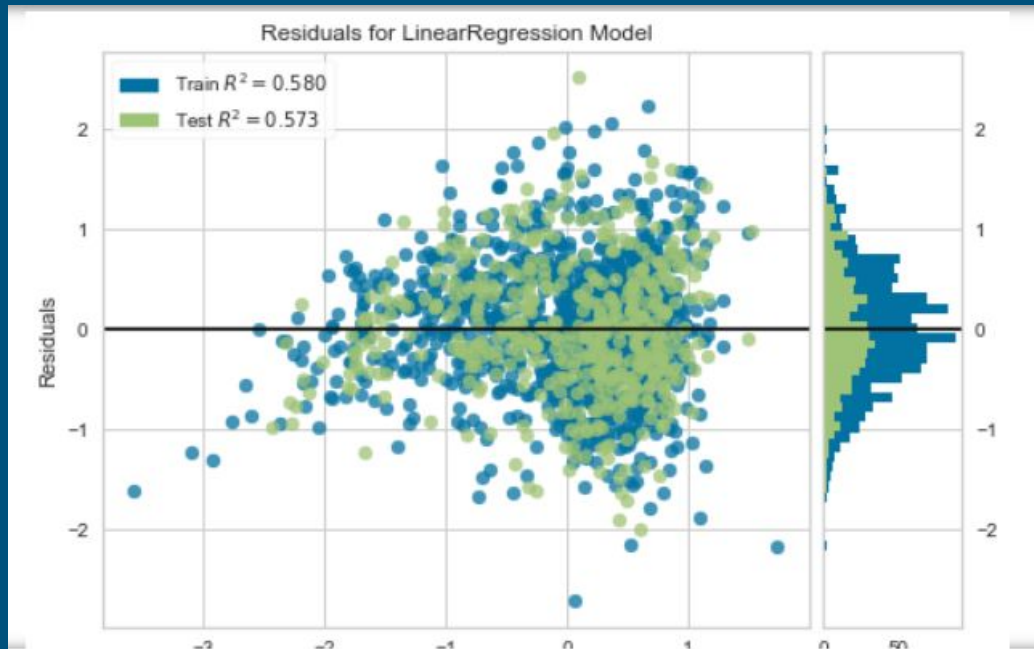
print("Training data score using ridge:", np.max(model_scores))

bestAlphaIndex = np.argmax(model_scores)
ridge_model.alpha = alpha_space[bestAlphaIndex]
ridge_model.fit(x_train, y_train)
print("Testing data score using lasso: ", ridge_model.score(x_test, y_test))
```

Training data score using ridge: 0.4596214850148262  
Testing data score using lasso: 0.482831796093244


# Making some checks

- I plotted the residuals to make sure everything looks okay
- There is a slight curve to the residuals
- Additionally, the data points show a trumpet like pattern
- However, points do appear to be normally distributed




# Making some checks cont.

- The trumpet shape indicates that the model might need additional predictor variables to explain the levels of ozone
- These additional variables probably include the hydrocarbons mentioned earlier, which aren't in the EPA dataset



# Making some check cont.

- Unfortunately, high quality, consistent air monitoring data is difficult to come by, so I'm going to keep using this data with the acknowledgement that this provides an incomplete picture



# Increasing the models' performance

- Next, I tried XGBoost, an implementation of gradient boosted decision trees
- Did considerably better with a score of 0.5908792859536407

```
#trying to improve the model using boosting

import xgboost as xgb

data_dmatrix = xgb.DMatrix(data=variables,label=target)
xg_reg = xgb.XGBRegressor(objective ='reg:squarederror', colsample_bytree = 0.3, learning_rate = 0.1,
                           max_depth = 5, alpha = 10, n_estimators = 100)
xg_reg.fit(x_train,y_train)

xg_test_pred = xg_reg.predict(x_test)
xg_train_pred = xg_reg.predict(x_train)

rmse = np.sqrt(mean_squared_error(y_train, xg_train_pred))
print("RMSE: %f" % (rmse))

rmse = np.sqrt(mean_squared_error(y_test, xg_test_pred))
print("RMSE: %f" % (rmse))

print("Training data score:", xg_reg.score(x_train,y_train))
print("Testing data score:", xg_reg.score(x_test,y_test))
#It's improved...a little bit

RMSE: 0.605974
RMSE: 0.639625
Training data score: 0.6327957699745272
Testing data score: 0.5908792859536407
```

# XGBoost Cross Validation

- For cross validation, I set the metric to RMSE in the interim as a stand in to score(), which I was using previously
- RMSE of base XGBoost model on the testing data set was 0.639625
- Cross validation significantly improved this to 0.013637

```
#try to improve using cross validation
#first set up hyperparameters
params = {"objective": "reg:squarederror", 'colsample_bytree': 0.3, 'learning_rate': 0.1,
          'max_depth': 5, 'alpha': 10}

cv_results = xgb.cv(dtrain=data_dmatrix, params=params, nfold=5,
                    num_boost_round=50, early_stopping_rounds=10, metrics="rmse", as_pandas=True, seed=2020)

#get the best best RSME
print(cv_results.iloc[len(cv_results)-1])
```

train-rmse-mean	0.013637
train-rmse-std	0.000043
test-rmse-mean	0.013652
test-rmse-std	0.000520

Name: 49, dtype: float64



# XGBoost Hyperparameter Tuning

- To make the XGBoost cross validation model better, I try to find the optimal hyperparameters
- Modified code from <https://blog.cambridgespark.com/hyperparameter-tuning-in-xgboost-4ff9100a3b2f>
- Started off with selecting an optimal max\_depth from the range of 3 to 9, inclusive

```
#with thanks to https://blog.cambridgespark.com/hyperparameter-tuning-in-xgboost-4ff9100a3b2f
gridsearch_params = [
    (max_depth)
    for max_depth in range(3,10)
]

min_rmse = float("Inf")
best_params = None

for max_depth in gridsearch_params:
    # Update our parameters
    params['max_depth'] = max_depth
    # Run CV
    cv_results = xgb.cv(dtrain=data_dmatrix, params=params, nfold=5,
                        num_boost_round=50, early_stopping_rounds=10, metrics="rmse", as_pandas=True, seed=2020)
    # Update best MAE
    mean_rmse = cv_results['train-rmse-mean'].min()
    boost_rounds = cv_results['train-rmse-mean'].argmin()
    if mean_rmse < min_rmse:
        min_rmse = mean_rmse
        best_params = (max_depth)

params["max_depth"] = best_params
print(params["max_depth"])
```

# Hyperparameter tuning cont.

- Next was learning\_rate

```
min_rmse = float("Inf")
best_params = None
for learning_rate in [.4, .3, .2, .1, .05, .01, .005]:
    # We update our parameters
    params['learning_rate'] = learning_rate
    # Run and time CV
    cv_results = xgb.cv(dtrain=data_dmatrix, params=params, nfold=5,
                        num_boost_round=50, early_stopping_rounds=10, metrics="rmse", as_pandas=True, seed=2020)
    # Update best score
    mean_rmse = cv_results['test-rmse-mean'].min()
    boost_rounds = cv_results['test-rmse-mean'].argmin()
    if mean_rmse < min_rmse:
        min_rmse = mean_rmse
        best_params = learning_rate

params["learning_rate"] = best_params
print(params["learning_rate"])
```

0.3

# Hyperparameter tuning cont.

- Third was `colsample_bytree` from a range of 0.1 to 0.6, inclusive, with steps of 0.1

```
gridsearch_params = [  
    (colsample)  
    for colsample in [i/10. for i in range(1, 7)]  
]  
  
min_rmse = float("Inf")  
best_params = None  
# We start by the largest values and go down to the smallest  
for colsample in reversed(gridsearch_params):  
    # We update our parameters  
    params['colsample_bytree'] = colsample  
    # Run CV  
    cv_results = xgb.cv(dtrain=data_dmatrix, params=params, nfold=5,  
                        num_boost_round=50, early_stopping_rounds=10, metrics="rmse", as_pandas=True, seed=2020)  
    # Update best score  
    mean_rmse = cv_results['test-rmse-mean'].min()  
    boost_rounds = cv_results['test-rmse-mean'].argmin()  
    if mean_rmse < min_rmse:  
        min_rmse = mean_rmse  
        best_params = (colsample)  
  
params["colsample_bytree"] = best_params  
print(params["colsample_bytree"])
```

0.5

# Hyperparameter tuning cont.

- Finally was adjusting alpha, selecting an integer between 7 and 13, inclusive

```
gridsearch_params = [
    (alpha)
    for alpha in range(6, 13)
]

min_rmse = float("Inf")
best_params = None
# We start by the largest values and go down to the smallest
for alpha in gridsearch_params:
    # We update our parameters
    params['alpha'] = alpha
    # Run CV
    cv_results = xgb.cv(dtrain=data_dmatrix, params=params, nfold=5,
                        num_boost_round=50, early_stopping_rounds=10, metrics="rmse", as_pandas=True, seed=2020)
    # Update best score
    mean_rmse = cv_results['test-rmse-mean'].min()
    boost_rounds = cv_results['test-rmse-mean'].argmin()
    if mean_rmse < min_rmse:
        min_rmse = mean_rmse
        best_params = (alpha)

params["alpha"] = best_params
print(params["alpha"])
```

# Hyperparameter tuning cont.

- The optimal hyperparameters based on my ranges are:
  - Colsample\_bytree = 0.5
  - Learning\_rate = 0.3
  - Max\_depth = 3
  - Alpha = 6

```
#create another model with the updated parameters
xg_reg = xgb.XGBRegressor(objective = 'reg:squarederror', colsample_bytree = 0.5, learning_rate = 0.3,
                           max_depth = 3, alpha = 7, n_estimators = 100)
xg_reg.fit(x_train,y_train)

xg_test_pred = xg_reg.predict(x_test)
xg_train_pred = xg_reg.predict(x_train)

print("Training data score using optimized parameters: ", xg_reg.score(x_train, y_train))
print("Testing data score using optimized parameters: ", xg_reg.score(x_test, y_test))
```

```
Training data score using optimized parameters:  0.7343487340606429
Testing data score using optimized parameters:  0.6183643073747713
```

# Hyperparameter tuning results

- This yields a RMSE of 0.009815, which is much lower than what was previously achieved

```
print(params)

cv_results = xgb.cv(dtrain=data_dmatrix, params=params, nfold=5,
                    num_boost_round=50, early_stopping_rounds=10, metrics="rmse", as_pandas=True, seed=2020)

#Printing out final RMSE
print(cv_results.iloc[len(cv_results)-1])
```

```
{'objective': 'reg:squarederror', 'colsample_bytree': 0.5, 'learning_rate': 0.3, 'max_depth': 3, 'alpha': 6}
train-rmse-mean      0.009815
train-rmse-std        0.000059
test-rmse-mean        0.009896
test-rmse-std         0.000322
Name: 49, dtype: float64
```

# Hyperparameter tuning results cont.

- But what does this translate to in terms of score()?
  - A much improved 0.6232516218994382

```
#create another model with the updated parameters
xg_reg = xgb.XGBRegressor(objective = 'reg:squarederror', colsample_bytree = 0.5, learning_rate = 0.3,
                           max_depth = 3, alpha = 6, n_estimators = 100)
xg_reg.fit(x_train,y_train)

xg_test_pred = xg_reg.predict(x_test)
xg_train_pred = xg_reg.predict(x_train)

print("Training data score using optimized parameters: ", xg_reg.score(x_train, y_train))
print("Testing data score using optimized parameters: ", xg_reg.score(x_test, y_test))
```

```
Training data score using optimized parameters: 0.7418738374233047
Testing data score using optimized parameters: 0.6232516218994382
```

# Which model is the best?

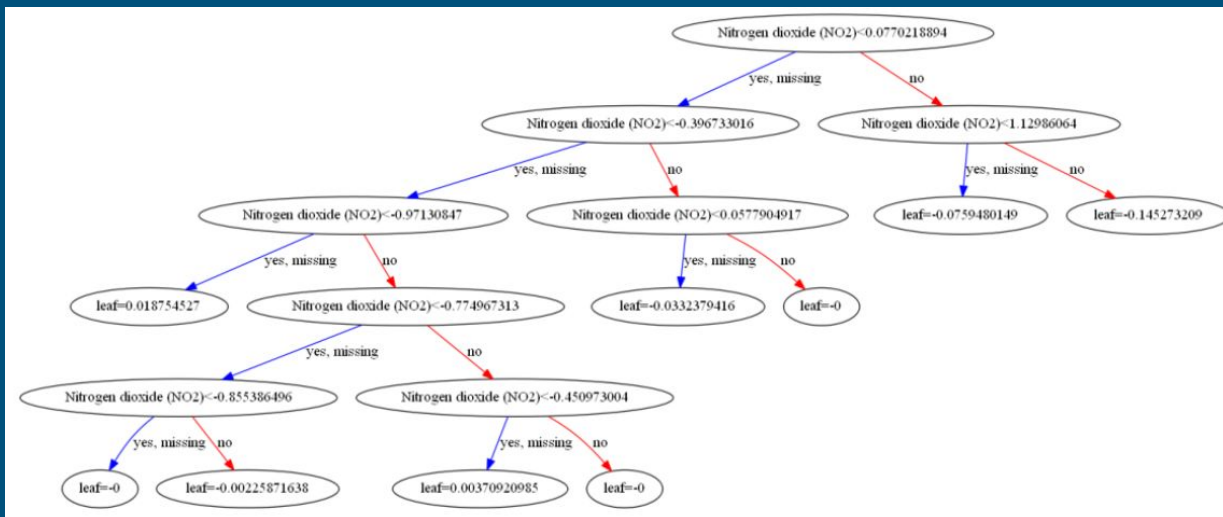
---

- The XGBoost model with optimized hyperparameters
- Still only ~60% accuracy, but that's okay considering the lack of inclusion of other relevant predictors
  - An accuracy of 60% is still a marked improvement over the base linear regression model



# Visualizing the best model's tree

- Here we can see the decision tree the final XGBoost model uses



# Ranking the importance of predictors

- And here's the importance of each predictor

